

# Data Motion Network and Optimization

SDx 2018.2



# Objectives

- After completing this module, you will be able to:
  - >> List various interface ports available to move data between the processor and user-generated IP in PL
  - >> Describe at the hardware level how caches interact with the accelerator coherency port (ACP)
  - >> List available functions to perform data exchanges and synchronize events
  - >> Identify some of the techniques used for hardware optimization

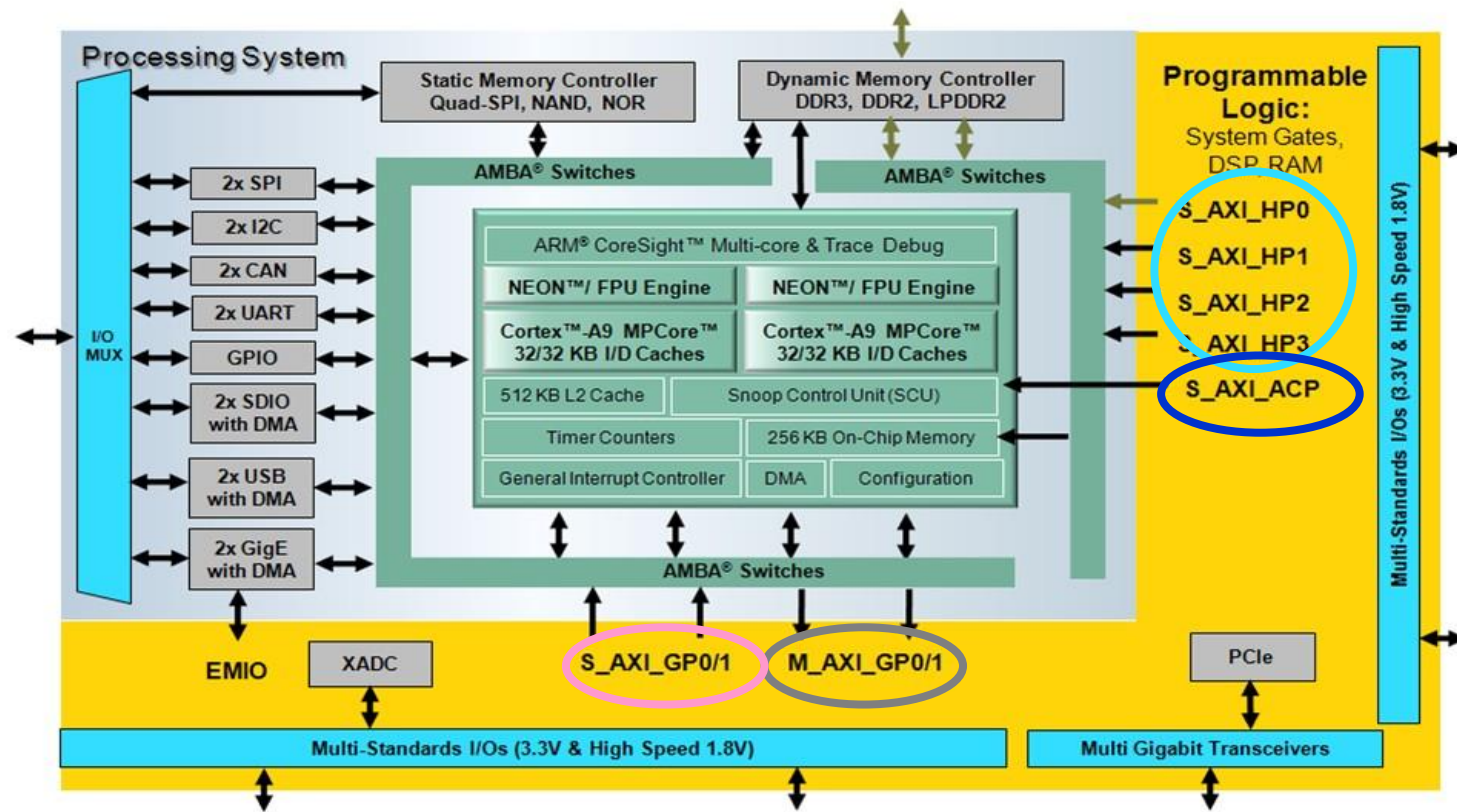
# Outline

- > Architectural Support for Data Motion
- > Data Motion
- > Optimization
- > Summary
- > Lab2 Intro

# PS-PL Interface AXI Ports

> The AMBA AXI ports of the PS-PL interface provide the primary mechanism for the flow of data between the PS and PL

- >> Two general-purpose master ports
- >> Two general-purpose slave ports
- >> Four high-performance slave ports
- >> One accelerator coherency port (ACP) slave port



# General-Purpose Master Ports

## > Two identical 32-bit AXI master ports

- >> M\_AXI\_GP0
- >> M\_AXI\_GP1

## > Attaches to slave AXI ports in programmable logic via PL AXI interconnect

- >> PL slave
  - Your peripheral built in programmable logic
  - IP Integrator interface or other third-party-based IP

## > Mostly used for CPU and I/O peripheral block data movement to programmable logic

## > Why two ports?


- >> Each port, having its own 1 GB space, is capable of driving a number of peripherals using an AXI switch
- >> Multiple ports enable the designer to distribute bandwidth

## > PS masters can be

- >> Cortex-A9 processors via the L2 cache controller
- >> USB, Ethernet, SD/SDIO controllers
- >> DMAC
- >> Debug access port

# General-Purpose Slave Ports

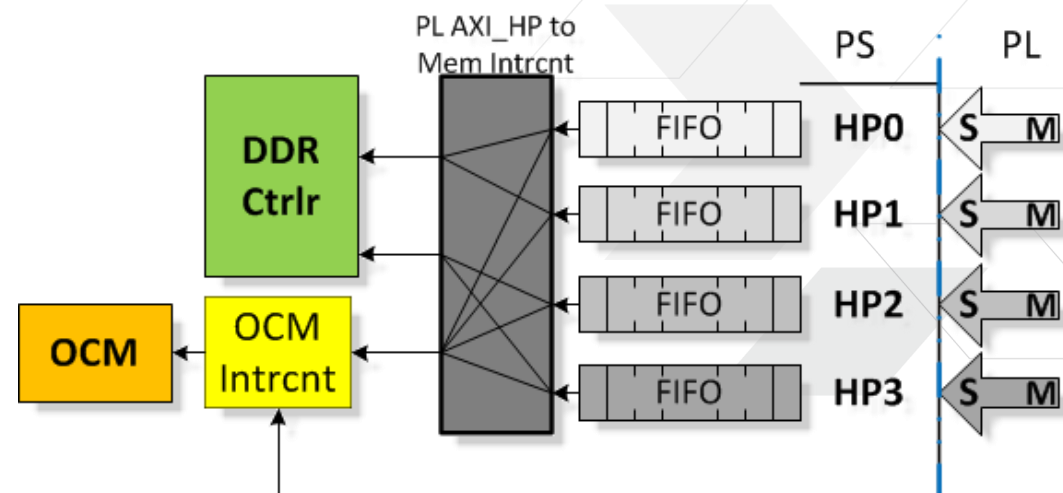
- > **Two identical 32-bit AXI slave ports**
  - >> S\_AXI\_GP0
  - >> S\_AXI\_GP1
- > **Provides initiator access from PL master to PS target**
- > **Attaches to master AXI port in PL via AXI interconnect**
  - >> PL master
    - MicroBlaze processor
    - Your master built in programmable logic
    - Other third-party-based IP
- > **PS slaves can be**
  - >> DDRx controller
  - >> On-chip memory (OCM)
  - >> IOP peripheral



# High-Performance Slave Ports

- > **Asynchronous crossing between the PL clock domain and PS clock domain**
  - >> S\_AXI\_HP0, S\_AXI\_HP1
  - >> S\_AXI\_HP2, S\_AXI\_HP3
- > **Provides initiator access from PL master to PS memory**
- > **Provides low latency access to DDR and OCM**
- > **Attaches to master AXI port in PL**
  - >> DMA controller
  - >> MicroBlaze processor
  - >> Custom master built in programmable logic
  - >> Other third-party based IP

- > **QoS supported from the programmable logic ports**
- > **FIFOs smooth out *long-latency* events**
  - >> 1 KB (128 by 64 bit) data FIFOs
  - >> Both read and write paths



# Accelerator Coherency Port (ACP) Slave Port

## > One accelerator coherence port (ACP)

- >> S\_AXI\_ACP
- >> High-performance, coprocessor interface to PL
  - Accelerators gain access to the CPU cache hierarchy
    - Enables PL to participate in the coherency among the caches and DDRx memory
  - Supported in hardware; no software needed
- >> No drivers required for the ACP; however, the IP connected to the ACP may require drivers

## > Provides initiator access from PL master to PS target

- >> Via L1 and L2 cache

## > Attaches to master AXI port in PL

- >> Typically a PL-based co-processor/accelerator
- >> Also could be
  - MicroBlaze processor
  - Custom master built in programmable logic
  - Other third-party-based IP



# Why Connection to the SCU is Important for Accelerators

- > **ACP bolts directly into the snoop control unit (SCU)**
  - >> ACP is a PS slave AXI port (that is, accelerator is a master)
  - >> Enables the accelerator on the ACP to write directly into the L1 and L2 caches
    - And indirectly to the DDR/OCM
  - >> Data movement is limited by the size of the cache target
    - Typically L2 is the target due to its larger size
    - Cache misses result in moving data to the DDR
    - Frequent misses indicate that the HP port (not the ACP port) should have been used
- > **When the ACP is not used, there is no non-standard use of the caches**

# System-Level Considerations that Determine Coding Techniques

- > **Which PS interfaces to use for the accelerator?**
  - >> AXI\_HP
  - >> AXI\_ACP
  - >> AXI\_GP (AXI\_lite for control)
- > **Where does the input to the accelerator connect from?**
  - >> From PS/PS DDR interface
  - >> MIG
  - >> From an external interface (like video)
- > **Where does the output from the accelerator connect to?**
  - >> To PS
  - >> To an external interface
- > **What is the optimal data movement mechanism between the PS and accelerator?**
  - >> DMA driven
    - Simple
    - Scatter gather
  - >> Software driven
  - >> How to manage cache coherency?
    - Software
    - SCU
- > **What is the optimal signaling between PS and the accelerator?**
  - >> Interrupt driven
  - >> Event interface
  - >> Polled

# Preferred Interfaces for Hardware Accelerators

- > **AXI\_HP DMA or AXI\_ACP DMA give the highest performance**
  - >> Both DDR and OCM (4x64K) accessible, lower latency through HP ports
- > **AXI\_ACP offers further performance when cache coherent data transfers are implemented to/from the accelerator**
- > **When is ACP more suited?**
  - >> Not suited for instruction-level acceleration; tightly coupled processor registers perform best
  - >> Not suited for large data sets; this can cause cache thrashing and degrade performance
  - >> Suited for small to moderate size data sets that can fit into L1 (<32K) or L2 (<256K) caches

# Comparing Data Movement Methods (1)

Method	Benefits	Drawbacks	Suggested Uses	Estimated Throughput
CPU Programmed I/O	<ul style="list-style-type: none"><li>• Simple software</li><li>• Fewest PL resources</li><li>• Simple PL slaves</li></ul>	<ul style="list-style-type: none"><li>• Lowest throughput</li></ul>	<ul style="list-style-type: none"><li>• Control functions</li></ul>	<25 MB/s
PS DMAC	<ul style="list-style-type: none"><li>• Fewest PL resources</li><li>• Medium throughput</li><li>• Multiple channels</li><li>• Simple PL slaves</li></ul>	<ul style="list-style-type: none"><li>• Somewhat complex DMA programming</li></ul>	<ul style="list-style-type: none"><li>• Limited PL resource DMAs</li></ul>	600 MB/s
PL AXI_HP DMA	<ul style="list-style-type: none"><li>• Highest throughput</li><li>• Multiple interfaces</li><li>• Command/data FIFOs</li></ul>	<ul style="list-style-type: none"><li>• OCM/DDR access only</li><li>• More complex PL master design</li></ul>	<ul style="list-style-type: none"><li>• High-performance DMA for large datasets</li></ul>	1200 MB/s (per interface)

# Comparing Data Movement Methods (2)

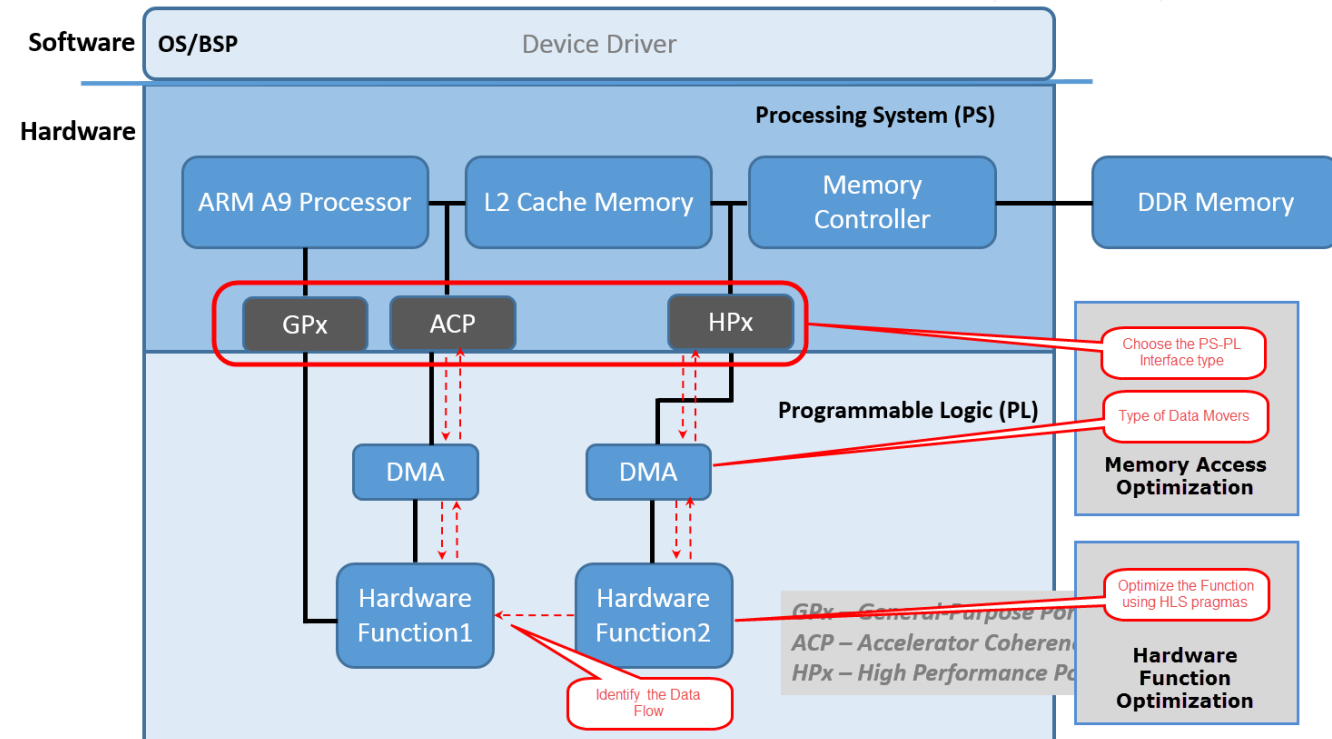
Method	Benefits	Drawbacks	Suggested Uses	Estimated Throughput
PL AXI_ACP DMA	<ul style="list-style-type: none"><li>• Highest throughput</li><li>• Lowest latency</li><li>• Optional cache coherency</li></ul>	<ul style="list-style-type: none"><li>• Large burst might cause cache thrashing</li><li>• Shared CPU interconnect bandwidth</li><li>• More complex PL master design</li></ul>	<ul style="list-style-type: none"><li>• High-performance DMA for smaller, Coherent datasets</li><li>• Medium granularity CPU offload</li></ul>	1200 MB/s
PL AXI_GP DMA	<ul style="list-style-type: none"><li>• Medium throughput</li></ul>	<ul style="list-style-type: none"><li>• More complex PL master design</li></ul>	<ul style="list-style-type: none"><li>• PL to PS control functions</li><li>• PS I/O peripheral access</li></ul>	600 MB/s

# Data Motion



# Data Movement in Zynq

- > Data may move between master/slave IP located in PL and memories located internal to the processor or external to the chip using the nine ports described in the previous section
- > The decision is made by the SDSoC compiler after analyzing the software or following user's directives



# How SDSoC Compiler Maps Programs to HW/SW?

- > **User specifies a program, a platform and a set of functions to map to hardware**
  - >> This defines the SDSoC-generated hardware / software interface
  - >> Program may also call other software functions provided by the platform, e.g., “write LED”, which are not targeted to hardware
- > **The SDSoC system compiler analyzes the program and maps it into a hardware / software system**
  - >> Preserves program semantics, including calls to hardware functions
  - >> Hardware functions can run concurrently with well-defined synchronization
- > **Hardware function calls define “connections” implemented using...**
  - >> Program properties: data flow between calls to hardware functions, memory allocation,...
  - >> Function argument properties: payload size, hardware interfaces,...
  - >> Function properties: memory access patterns, implementation latency,...



# How SDSoC Compiler Maps Programs to HW/SW? (2)

## > Program source code

```
bool mmultadd_test(float *A, float *B, float *C, float *Ds, float *D)
{
    float tmp1[A_NROWS * A_NCOLS], tmp2[A_NROWS * A_NCOLS];

    for (int i = 0; i < NUM_TESTS; i++) {
        mmultadd_init(A, B, C, Ds, D);

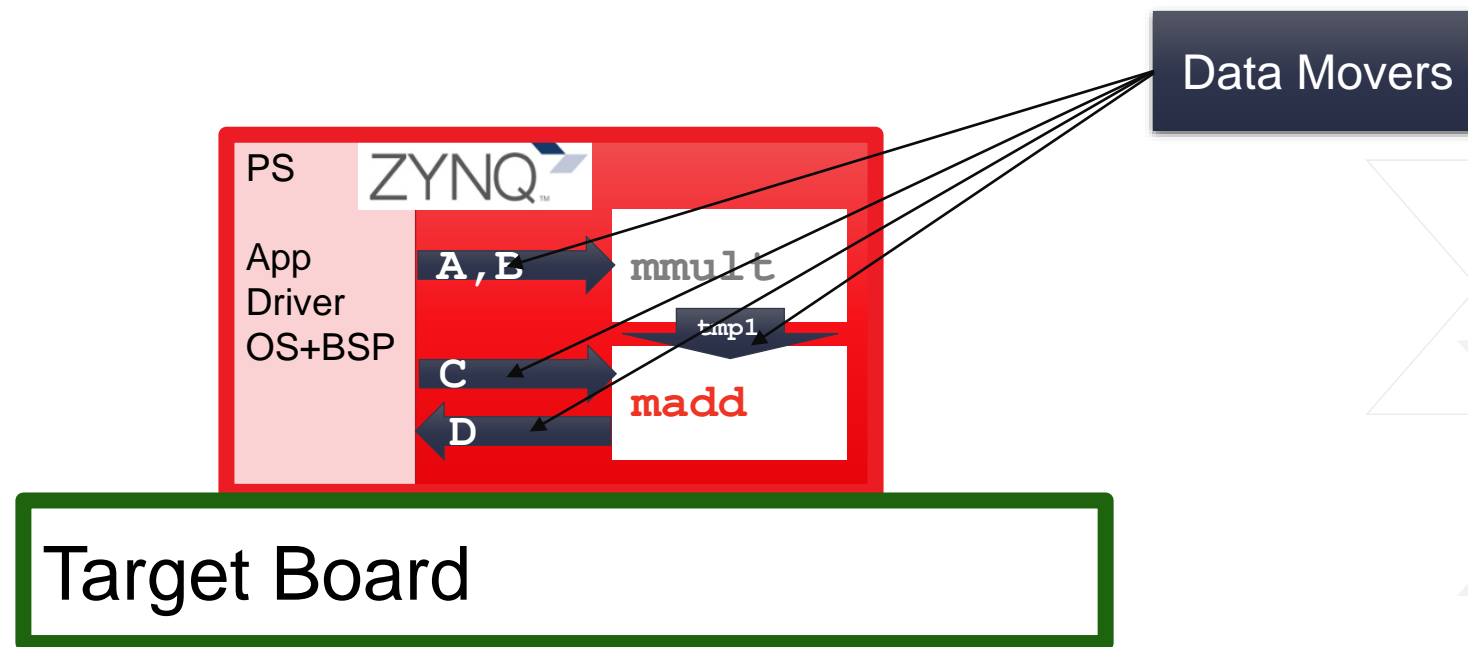
        mmult(A, B, tmp1);
        madd(tmp1, C, D);

        mmult_sw(A, B, tmp2);
        madd_sw(tmp2, C, Ds);

        if (!mmult_result_check(D, Ds))
            return false;
    }
    return true;
}
```

# How SDSoC Compiler Maps Programs to HW/SW? (3)

## ➤ Structure of generated hardware system



# How SDSoC Compiler Maps Programs to HW/SW? (4)

## ➤ Structure of generated software

```
bool mmultadd_test(float *A, float *B, float *C, float *Ds, float *D)
{
    std::cout << "Testing mmult .." << std::endl;

    float tmp1[A_NROWS * A_NCOLS], tmp2[A_NROWS * A_NCOLS];

    for (int i = 0; i < NUM_TESTS; i++) {
        mmultadd_init(A, B, C, Ds, D);

        _p0_mmult_0(A, B, tmp1);
        //std::cout << "tmp1[0] = " << tmp1[0] << std::endl;
        _p0_madd_0(tmp1, C, D);

        mmult_golden(A, B, tmp2);
        madd_golden(tmp2, C, Ds);

        if (!mmult_result_check(D, Ds))
            return false;
    }

    return true;
}
```

```
void _p0_mmult_0(float in_A[1024], float in_B[1024], float out_C[1024])
{
    switch_to_next_partition(0);
    int start_seq[3];
    start_seq[0] = 0x00000000;
    start_seq[1] = 0x00010000;
    start_seq[2] = 0x00020000;
    cf_request_handle_t _p0_swinst_mmult_0_cmd;
    cf_send_i(&(_p0_swinst_mmult_0.cmd_mmult),
              start_seq, 3*sizeof(int), &_p0_swinst_mmult_0_cmd);
    cf_wait(_p0_swinst_mmult_0_cmd);
```

Control transfer

```
void _p0_madd_0(float A[1024], float B[1024], float C[1024])
{
    switch_to_next_partition(0);
    int start_seq[3];
    start_seq[0] = 0x00000003;
    start_seq[1] = 0x00010001;
    start_seq[2] = 0x00020000;
    cf_request_handle_t _p0_swinst_madd_0_cmd;
    cf_send_i(&(_p0_swinst_madd_0.cmd_madd), start_seq, 3*sizeof(int),
              &_p0_swinst_madd_0_cmd);
    cf_wait(_p0_swinst_madd_0_cmd);

    cf_send_i(&(_p0_swinst_madd_0.B_PORTA), B, 1024 * 4, &_p0_request_0);
    cf_receive_i(&(_p0_swinst_madd_0.C_PORTA), C, 1024 * 4,
                 &_p0_madd_0_num_C_PORTA, &_p0_request_1);
    cf_wait(_p0_request_0);
    cf_wait(_p0_request_1);
    cf_wait(_p0_request_2);
    cf_wait(_p0_request_3);
}
```

Control transfer

Data transfers

```
mult_0.in_A), in_A, 1024 * 4, &_p0_request_2);
mult_0.in_B), in_B, 1024 * 4, &_p0_request_3);
```

Data transfers

# Optimization



# System Level Optimization Using SDSoC

- > Program connectivity and implementation
- > Datamovers
- > Memory models
- > Platform port selection
- > Direct connections
- > Hardware function interface



# Program Connectivity


```
bool mmultadd_test(float *A, float *B, float *C, float *Ds, float *D)
{
    float tmp1[A_NROWS * A_NCOLS], tmp2[A_NROWS * A_NCOLS];

    for (int i = 0; i < NUM_TESTS; i++) {
        mmultadd_init(A, B, C, Ds, D);

        mmult(A, B, tmp1);
        //std::cout << "tmp1[0] = " << tmp1[0] << std::endl;
        madd(tmp1, C, D);

        mmult_golden(A, B, tmp2);
        madd_golden(tmp2, C, Ds);

        if (!mmult_result_check(D, Ds))
            return false;
    }
    return true;
}
```



# Program Connectivity (2)


```
bool mmultadd_test(float *A, float *B, float *C, float *Ds, float *D)
{
    float tmp1[A_NROWS * A_NCOLS], tmp2[A_NROWS * A_NCOLS];

    for (int i = 0; i < NUM_TESTS; i++) {
        mmultadd_init(A, B, C, Ds, D);

        mmult(A, B, tmp1);
        std::cout << "tmp1[0] = " << tmp1[0] << std::endl;
        madd(tmp1, C, D);

        mmult_golden(A, B, tmp2);
        madd_golden(tmp2, C, Ds);

        if (!mmult_result_check(D, Ds))
            return false;
    }
    return true;
}
```



> Simple program changes can cause significant system changes!

# Program Connectivity Implementation

- > Every connection with data movement between the software program and a hardware function requires a *datamover*
- > Program connections can have multiple possible implementations
  - >> SDSoC compiler infers datamovers and system ports
  - >> Datamover and system ports can be overridden with pragma
- > Datamover inference is based on payload size, payload memory attributes, and hardware function argument access pattern
  - >> `#pragma SDS data data_mover(arg:<DM_TYPE>) // to override`
- > Note: `#pragma SDS` is always treated as a rule, not a hint
  - >> Not all errors will be caught, so care is required



# Datamovers

## > Every datamover has two components

- >> Hardware IP, e.g., axi\_dma, a hardware function or a CPU
- >> OS-specific software library function (userspace for Linux)

Datamover	IP	IP port types	Payload (bytes)	Phys. memory contiguity
axilite	ps7	register, axilite		either
axi_dma_simple	axi_dma	bram,ap_fifo, axis	< 8M	contiguous
axi_dma_sg	axi_dma	bram,ap_fifo, axis		either
axi_dma_2d	axi_dma	bram		contiguous
axi_fifo	axi_fifo_mm_s	bram,ap_fifo, axis	≤ 300	non-contiguous
zero_copy	HW fcn IP	m_aximm		contiguous

# Memory Models

## > **sds\_lib.h** functions to “declare” memory properties

```
>> sds_alloc(size_t size); // guaranteed
contiguity
>> sds_mmap(void *paddr, size_t size, void *vaddr); // assumed contiguity
>> sds_register_dmabuf(void *vaddr, int fd); // assumed
contiguity
```

## > **#pragma SDS data copy (A[offset:array\_size])**

```
>> Copy in, copy out semantics (e.g., axi_dma, axi_fifo)
>> Also used to specify variable and non-default transfer sizes
```

## > **#pragma SDS zero\_copy (A[offset:array\_size])**

```
>> Shared memory semantics (hardware function m_aximm only)
```

## > **#pragma SDS data access\_pattern (A:SEQUENTIAL|RANDOM)**

```
>> Only valid for “copy” semantics
>> SDSoC generates FIFO or BRAM hardware interface accordingly
```

# Platform Port Selection

- > **Compiler infers ACP or HP (AFI) ports based on memory attributes and other criteria**

- >> `#pragma SDS data sys_port (A:ACP|AFI|MIG)` to override inference rules
  - A must be one of the formal arguments of the function

- > **Memory is by default assumed CACHEABLE, i.e., compiler maintains cache coherency between CPU and hardware function**

- >> Cache flush before transferring data to hardware function
- >> Cache invalidate before transferring data from hardware function

# Direct Hardware Connections

- > Compiler maps data movement between hardware functions onto direct connections when possible

- > Direct connections implemented entirely in hardware via AXI streams

Port IF	BRAM	FIFO	AXIS
BRAM	Y	Y	64-bit TDATA
FIFO	Y	Y	64-bit TDATA
AXIS	64-bit TDATA	64-bit TDATA	Equal TDATA width

- > User must guarantee sufficient buffering

>> `#pragma SDS data buffer_depth(A:<buffer_depth>)`

Interface	Type	Depth	Array Size	Data width	Default
BRAM	multi-buffer	1:4	2:16384	8:16:32:64	1
FIFO	buffer	2:(16384/array_size)	2:16384	8:16:32:64	1024

# Hardware Function Call Sequencing

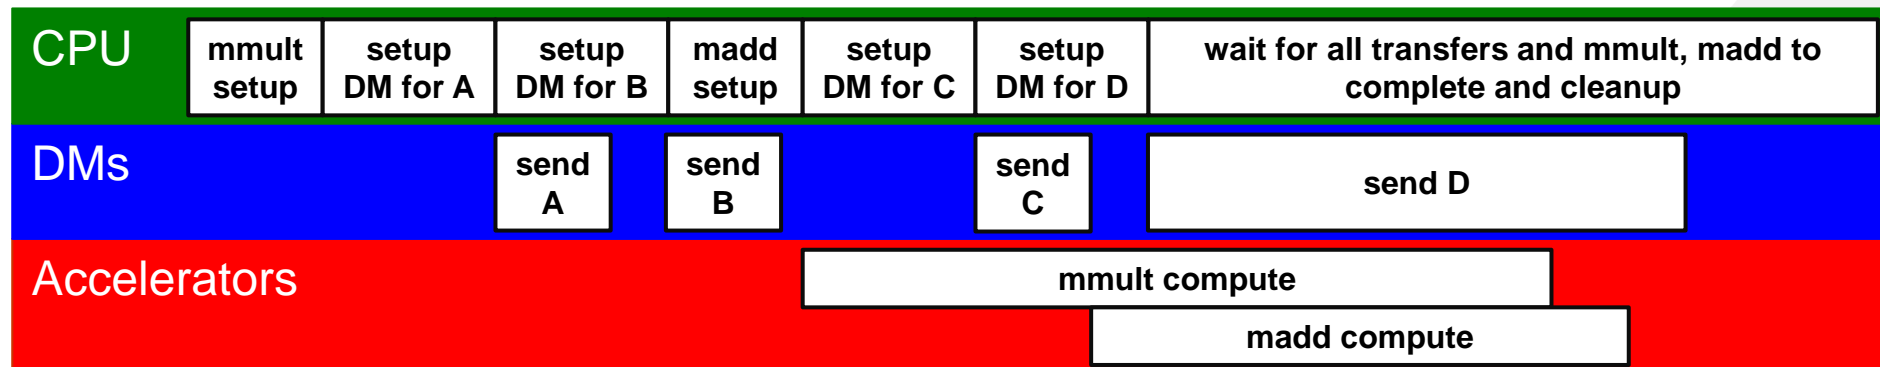
```
bool mmultadd_test(float *A, float *B, float *C, float *Ds, float *D)
{
    float tmp1[A_NROWS * A_NCOLS], tmp2[A_NROWS * A_NCOLS];

    for (int i = 0; i < NUM_TESTS; i++) {
        mmultadd_init(A, B, C, Ds, D);

        mmult(A, B, tmp1);
        //std::cout << "tmp1[0] = " << tmp1[0] << std::endl;
        madd(tmp1, C, D);

        mmult_golden(A, B, tmp2);
        madd_golden(tmp2, C, Ds);

        if (!mmult_result_check(D, Ds))
            return false;
    }
    return true;
}
```



# Hardware Function Call Interface

- > SDSoC rewrites calls to hardware functions and replaces function by a stub to transfer data and control

```
bool mmultadd_test(float *A, float *B, float *C, float *Ds, float *D)
{
    std::cout << "Testing mmult .." << std::endl;

    float tmp1[A_NROWS * A_NCOLS], tmp2[A_NROWS * A_NCOLS];

    for (int i = 0; i < NUM_TESTS; i++) {
        mmultadd_init(A, B, C, Ds, D);

        _p0_mmult_0(A, B, tmp1);
        //std::cout << "tmp1[0] = " << tmp1[0] << std::endl;
        _p0_madd_0(tmp1, C, D);

        mmult_golden(A, B, tmp2);
        madd_golden(tmp2, C, Ds);

        if (!mmult_result_check(D, Ds))
            return false;
    }

    return true;
}
```

In \_sds/swstubs directory

```
void _p0_mmult_0(float in_A[1024], float in_B[1024], float out_C[1024])
{
    switch_to_next_partition(0);
    int start_seq[3];
    start_seq[0] = 0x00000000;
    start_seq[1] = 0x00010000;
    start_seq[2] = 0x00020000;
    cf_request_handle_t _p0_swinst_mmult_0_cmd;
    cf_send_i(&(_p0_swinst_mmult_0.cmd_mmult),
              start_seq, 3*sizeof(int), &_p0_swinst_mmult_0_cmd);
    cf_wait(_p0_swinst_mmult_0_cmd);

    cf_send_i(&(_p0_swinst_mmult_0.in_A), in_A, 1024 * 4, &_p0_request_2);
    cf_send_i(&(_p0_swinst_mmult_0.in_B), in_B, 1024 * 4, &_p0_request_3);
}
```

Control transfer

Data transfers

# Hardware Function Call Interface (2)

- > `cf_wait()` synchronization barriers automatically inserted at the end of hardware function pipeline

```
bool mmultadd_test(float *A, float *B, float *C, float *Ds, float *D)
{
    std::cout << "Testing mmult .." << std::endl;

    float tmp1[A_NROWS * A_NCOLS], tmp2[A_NROWS * A_NCOLS];

    for (int i = 0; i < NUM_TESTS; i++) {
        mmultadd_init(A, B, C, Ds, D);

        _p0_mmult_0(A, B, tmp1);
        //std::cout << "tmp1[0] = " << tmp1[0] << std::endl;
        _p0_madd_0(tmp1, C, D);
    }
}
```

```
void _p0_mmult_0(float in_A[1024], float in_B[1024], float out_C[1024])
{
    switch_to_next_partition(0);
    int start_seq[3];
    start_seq[0] = 0x00000000;
    start_seq[1] = 0x00010000;
    start_seq[2] = 0x00020000;
    cf_request_handle_t _p0_swinst_mmult_0_cmd;
    cf_send_i(&(_p0_swinst_mmult_0.cmd_mmult),
              start_seq, 3*sizeof(int), &_p0_swinst_mmult_0_cmd);
    cf_wait(_p0_swinst_mmult_0_cmd);

    cf_send_i(&(_p0_swinst_mmult_0.in_A), in_A, 1024 * 4, &_p0_request_2);
    cf_send_i(&(_p0_swinst_mmult_0.in_B), in_B, 1024 * 4, &_p0_request_3);
}
```

```
void _p0_madd_0(float A[1024], float B[1024], float C[1024])
{
    switch_to_next_partition(0);
    int start_seq[3];
    start_seq[0] = 0x00000003;
    start_seq[1] = 0x00010001;
    start_seq[2] = 0x00020000;
    cf_request_handle_t _p0_swinst_madd_0_cmd;
    cf_send_i(&(_p0_swinst_madd_0.cmd_madd), start_seq, 3*sizeof(int),
              &_p0_swinst_madd_0_cmd);
    cf_wait(_p0_swinst_madd_0_cmd);

    cf_send_i(&(_p0_swinst_madd_0.B_PORTA), B, 1024 * 4, &_p0_request_0);
    cf_receive_i(&(_p0_swinst_madd_0.C_PORTA), C, 1024 * 4,
                 &_p0_madd_0_num_C_PORTA, &_p0_request_1);

    cf_wait(_p0_request_0);
    cf_wait(_p0_request_1);
    cf_wait(_p0_request_2);
    cf_wait(_p0_request_3);
}
```

← `cf_wait()` for madd

← `cf_wait()` for mmult

# Summary





# Summary

- > **High-performance AXI port connects the accelerator to DDR/OCM or SCU for data moving**
  - >> Internal routing and port select can provide low-latency paths between the accelerator and memory
- > **Caches provide both processors and accelerator with coherent data**
  - >> Caches can be locked and flushed as needed for the accelerator
- > **Processor provides clock sources to accelerators**
- > **Compiler rewrites calls to hardware functions and replaces function by a stub to transfer data and control**
- > **There are three functions available to exchange data and synchronize events**
  - >> `cf_send`
  - >> `cf_receive`
  - >> `cf_wait`

# Lab2 Intro



# Lab2 Intro

## > Introduction

- >> This lab guides you through the process of handling data movements between the software and hardware accelerators using various pragmas and functions

## > Objectives

- >> Use pragmas to select ACP or AFI ports for data transfer
- >> Use pragmas to select different data movers for your hardware function arguments
- >> Understand the use of `sds_alloc()` and `sds_free()` calls
- >> Understand the use of `malloc()` and `free()` calls
- >> Analyze built hardware

**Adaptable.**  
**Intelligent.**

