

# Using C-Callable libraries and creating multiple accelerators

SDx 2018.2



# Objectives

- > **After completing this module, you will be able to:**
  - >> List components of C-Callable functions
  - >> Describe the default behavior of the SDSoC development environment for handling multiple accelerators
  - >> Generate multiple accelerators for the same function
  - >> Use a pragma to generate multiple instances of an accelerator to override tool defaults
  - >> Explain the difference between blocking and non-blocking architectures as it relates to the SDSoC development environment

# Outline

- > C-Callable Accelerators
- > Multiple Accelerators
- > Blocking vs Non-Blocking calls
- > Summary

# SDSoC System Compilation and Linking

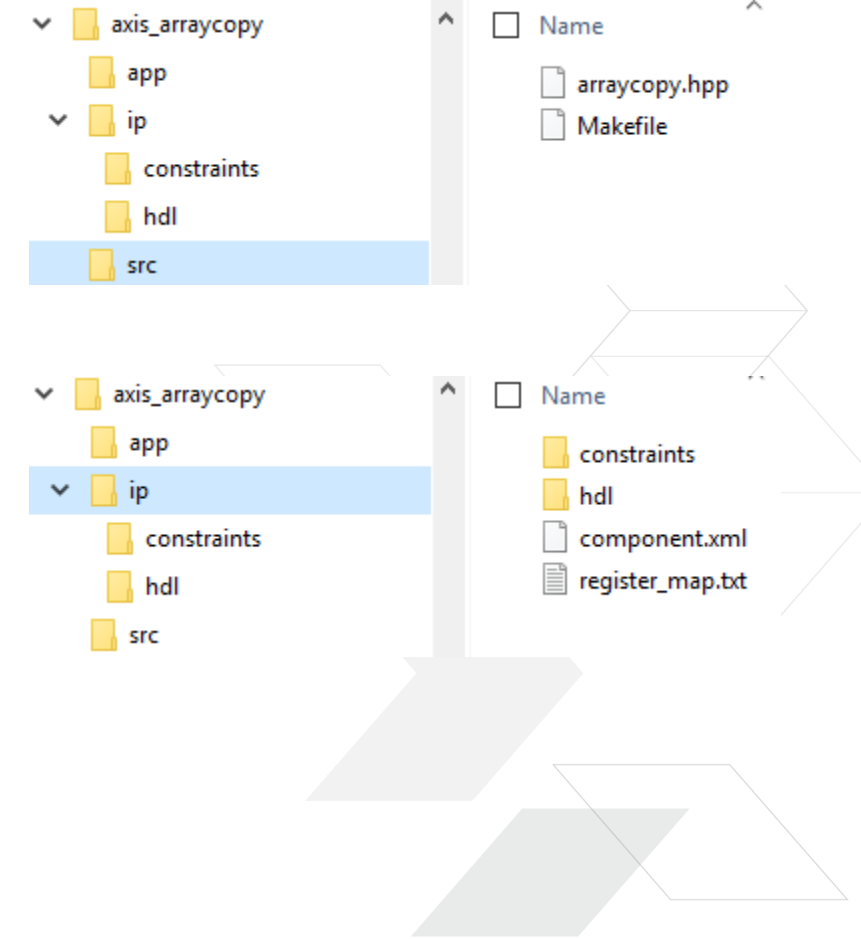
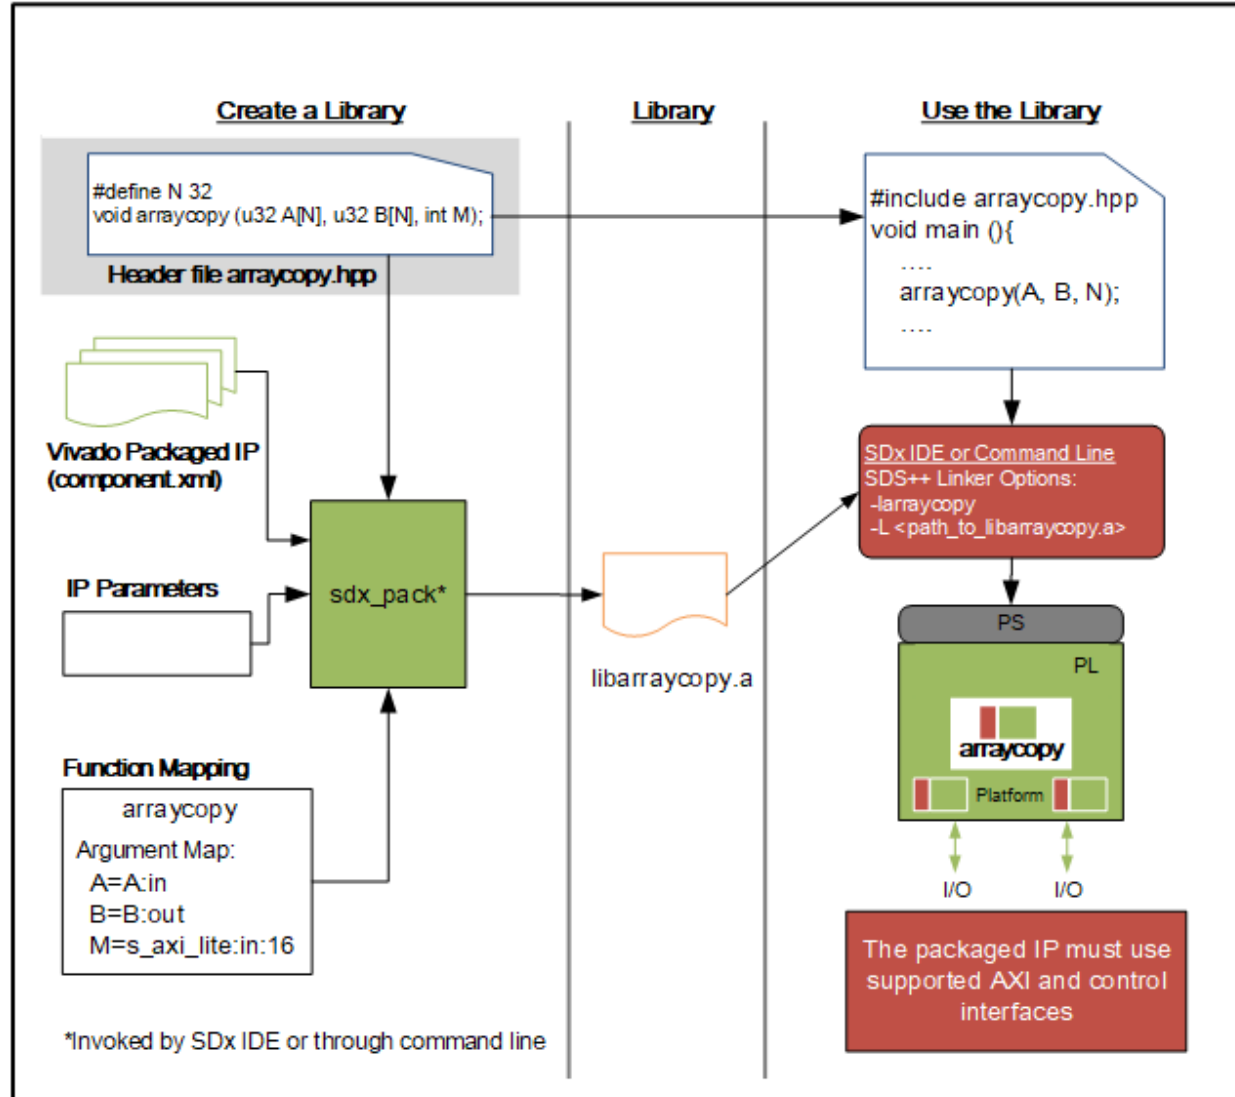
## > SDSoC tool compiler

- >> sdsc/sds++ compiler built upon Clang/LLVM frameworks, generates LLVRM (.s) files for each file, creating HLS interface pragmas for accelerators
- >> Creates a local HLS project for each accelerator and calls HLS
  - HLS returns an IP packaged by the Vivado Design Suite for each accelerator
  - HLS returns an interface XML file (database)
- >> Encapsulates everything needed for system linking into a standard .o file
- >> C-callable IP must have all of the above
  - User responsibility to provide various components
- >> SDSoC system linker: Takes .o files, analyzes entire program to generate accelerator network and software stubs

# What is C-Callable Library?

- > **C-Callable library consists of IP blocks written in VHDL or Verilog which can be called from user applications developed in SDSoC by statically linking it**
  - >> The IP blocks will be instantiated into the generated hardware system
- > **C-Callable library function consists of two components**
  - >> Header File
    - Function prototype
  - >> Static Library
    - Function definition
    - IP core
    - IP configuration parameters
    - Function argument mapping

# C-Callable Library Usage Flow



# C-Callable Function Components (1)

## > Header file

>> A library must declare function prototypes that map onto the IP block in a header file that can be #included in user application source files

>> Example:

```
// FILE: arraycopy.hpp

#ifndef _ARRAYCOPY_H_
#define _ARRAYCOPY_H_
#define N 32
typedef unsigned int u32;
#pragma SDS data copy(A[0:M])
#pragma SDS data copy(B[0:M])
void arraycopy(u32 A[N], u32 B[N], int M);
#endif
```

# C-Callable Function Components (2)

## > IP Core

- >> An HDL IP core for a C-callable library must be packaged using the Vivado tools
  - The packager creates a directory structure for the HDL and other source files, and an IP Definition file
  - The IP control register must exist at address offset 0x0
  - The `ap_start` signal initiates the IP execution, `ap_done` indicates IP task completion, and `ap_ready` indicates that the IP can be started
- >> Place the IP core either in the Vivado tools IP repository or in any other location
- >> Example

```
// bit 0 - ap_start (Read/Write/COH)
// bit 1 - ap_done (Read/COR)
// bit 2 - ap_idle (Read)
// bit 3 - ap_ready (Read)
// bit 7 - auto_restart (Read/Write)
// (COR = Clear on Read, COH = Clear on Handshake)
```



# C-Callable Function Components (3)

## > IP Configuration Parameters

- >> Most HDL IP cores are customizable at synthesis time
- >> Customization done through IP parameters
- >> SDSoC uses this information during the core instantiation in a generated system
- >> The information is captured in an XML file

– Example:

```
<?xml version="1.0" encoding="UTF-8"?>
<spirit:component xmlns:xilinx="http://www.xilinx.com" xmlns:spirit="http://www.spiritconsortium.org/XMLSchema/SPIRIT/1685-2009" xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">
  <spirit:vendor>xilinx.com</spirit:vendor>
  <spirit:library>rtl</spirit:library>
  <spirit:name>arraycopy</spirit:name>
  <spirit:version>1.0</spirit:version>
  <spirit:busInterfaces>
    <spirit:busInterface>
      <spirit:name>s_axi_lite</spirit:name>
      <spirit:busType spirit:vendor="xilinx.com" spirit:library="interface" spirit:name="aximm" spirit:version="1.0"/>
      <spirit:abstractionType spirit:vendor="xilinx.com" spirit:library="interface" spirit:name="aximm_rtl" spirit:version="1.0"/>
      <spirit:slave>
        <spirit:memoryMapRef spirit:memoryMapRef="s_axi_lite"/>
      </spirit:slave>
    </spirit:busInterface>
  </spirit:busInterfaces>
  <spirit:portMaps>
```

# C-Callable Function Components (6)

- > A utility called `sdx_pack` allows the creation of SDSoC libraries

`sdx_pack [arguments] [options]`

- >> **Example:**

```
sdx_pack -header arraycopy.hpp -lib libarraycopy.a \  
-func arraycopy -map A=A:in -map B=B:out -map M=s_axi_lite:in:16 -func-end \  
-ip ../ip/component.xml -control AXI=s_axi_lite:0 \  
-target-family zynq -target-cpu cortex-a9 -target-os standalone \  
-verbose
```

# How to Use C-Callable Library

- > Using a C-callable library is similar to using any software library
  - >> Use `#include` header files for the library in appropriate source files
  - >> Use the `sdscc -I<path>` option to compile your source from command line
  - >> Use C/C++ Build Settings->SDSCC Compiler->Directories (or SDS++ Compiler->Directories for C++ compilation) in GUI
  - >> Link the library using the `-L<path>` and `-l<lib>` options from the command line
  - >> Use C/C++ Build Settings > SDS++ Linker > Libraries in GUI
    - Use `lib<library_name>.a` as the library name and `<library_name>` with `-l` switch
    - For example, `libMylib.a` as the library name and `-lMylib` as the switch

# Multiple Accelerators



# Default Behaviors for Accelerator Replication

- > Multiple calls to a hardware function are usually mapped to a single accelerator
- > SDSoC tool will automatically generate multiple instances of an accelerator
- > Multiple instances will create parallelism and can improve system performance
- > If the output of an accelerator is an input of another accelerator, SDSoC tool will ***pipeline*** the accelerators

- >> Avoiding unnecessary access to main memory

- >> The compiler creates two instances of the hardware function as shown below

```
mmult_accel(tin1Buf, tin2Buf, toutBufHwInter);  
mmult_accel(toutBufHwInter, tin2Buf, toutBufHw);
```

# Forcing Multiple Instances of an Accelerator


- > **#pragma SDS resource(id) before function call**
  - >> Direct the compiler to create multiple instances of a hardware by inserting this pragma immediately preceding a call to the function
  - >> Directs SDSoC tool to return control after setting up all data transfers
  - >> Different ID for the same accelerator function results in different accelerator hardware instance
- > **#pragma SDS wait(id) must be used at suitable synchronization points**
- > **SDSoC tool may create multiple accelerators to improve performance**
  - >> Sometimes helpful for PL-to-PL connections

```
#pragma SDS resource(1)
mmult_accel(...); // instance 1
#pragma SDS resource(2)
mmult_accel(...); // instance 2
...
#pragma SDS wait(1)
#pragma SDS wait(2)
```

# No Estimate Support for Asynchronous Calls

- > **Generating multiple instances of an accelerator affects the macro-architecture**

- >> Is very difficult to provide accurate estimations with the current SDEstimate flow
  - Therefore, for the current version, there is no support for the performance estimation flow



# Task-Level Pragmas in the SDSoC Development Environment

## > Some additional SDS pragmas are necessary to implement task-level pipelining

>> `#pragma SDS async(ID)`

- Tells the compiler that the following accelerator call should be executed asynchronously
- ID specifies a unique ID for each use of `async()` with a specific accelerator
- Separate accelerator instance is generated for each unique ID

>> `#pragma SDS wait(ID)` or `function sds_wait(ID)`

- Allows a program to wait for completion of an asynchronously called accelerator
- Task completion is always triggered in the same order in which they were called
- Run-time queue maintains wait-handles associated with the arguments of each task

>> `#pragma SDS buffer_depth(Arg1:BufferDepth, Arg2:BufferDepth, ...)`

- Creates the multi-buffers required to pipeline
- Depth must be compatible with pipeline length

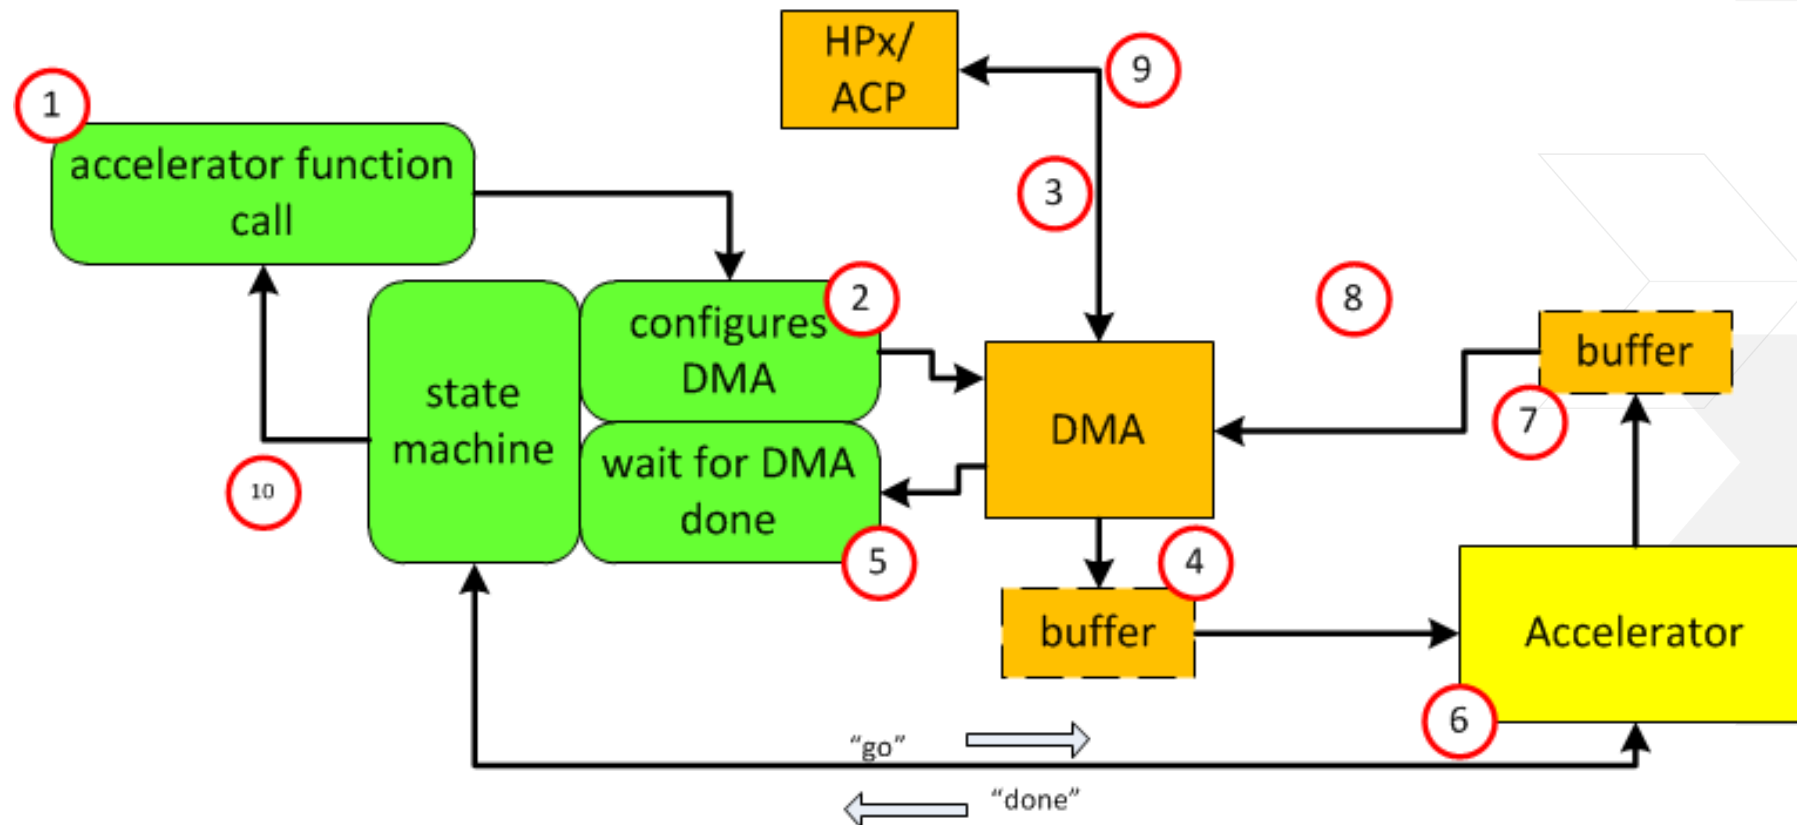


# Blocking vs Non-Blocking Calls




# An Accelerator Task

- > The *task* of calling an accelerator involves many steps in the background

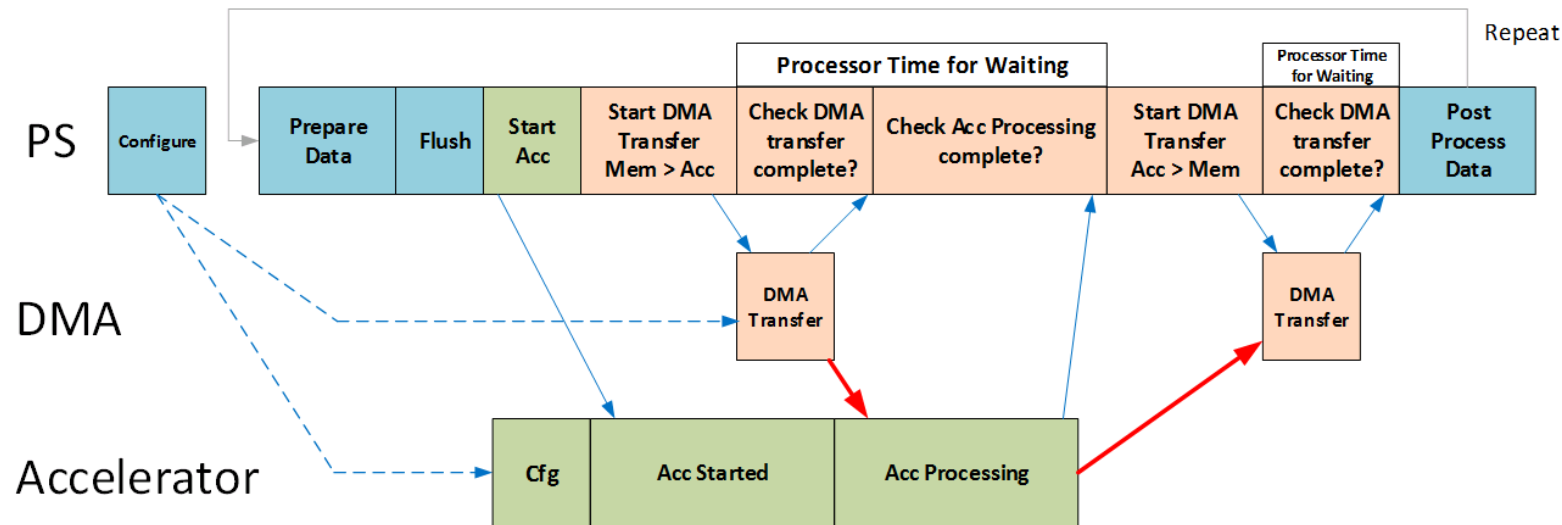


# Blocking Behavior

- > **Blocking behavior causes the processor to stall during accelerator task**
  - >> Sending input arguments
  - >> Processing of arguments in the accelerator
  - >> Returning results
- > **Benefit is reduced coding complexity (no synchronization issues)**
- > **Drawback is that processor sits idle (waste processing power)**
- > **Referred to as *sequential* in the SDSoC development environment**




# Accelerator I/O Connected to PS: Blocking



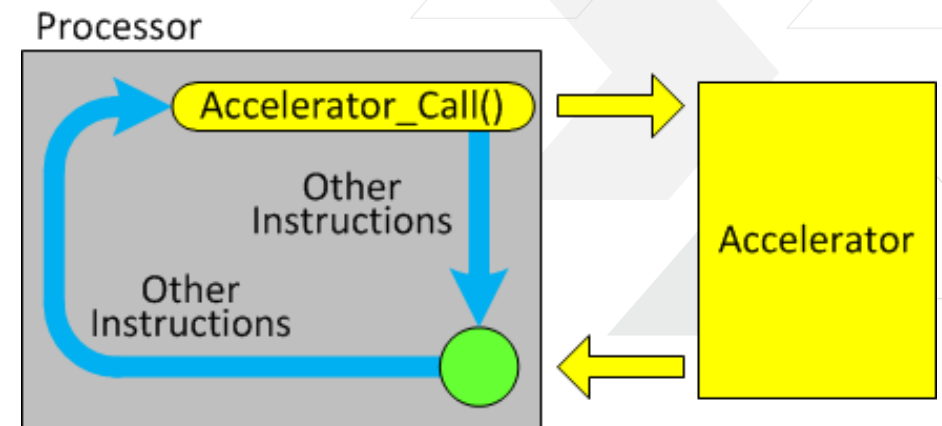
Flushing of Cache is required when the DMA can only access the DDR and not the Cache directly

Control Flow (AXI\_lite)  
Data Flow (AXI\_HP, AXI\_ACP)

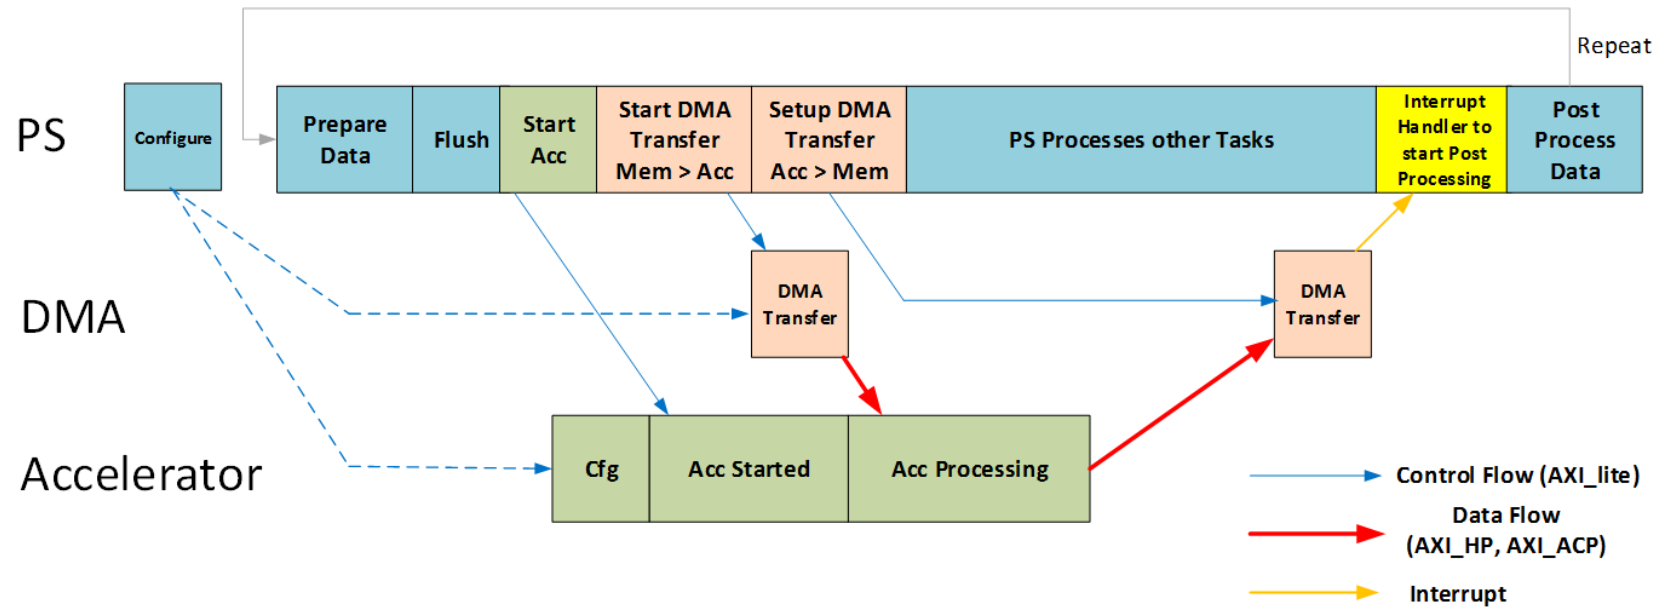


# Non-Blocking Behavior


- > True non-blocking (or just non-blocking) behavior allows the processor to continue operation after the accelerator call is made
- > Benefit is that processor can continue operation simultaneously with accelerator
- > Drawback is additional complexity (synchronization issues)
- > SDSoC refers to this as *asynchronous* calling
  - >> Only applies to accelerator calls



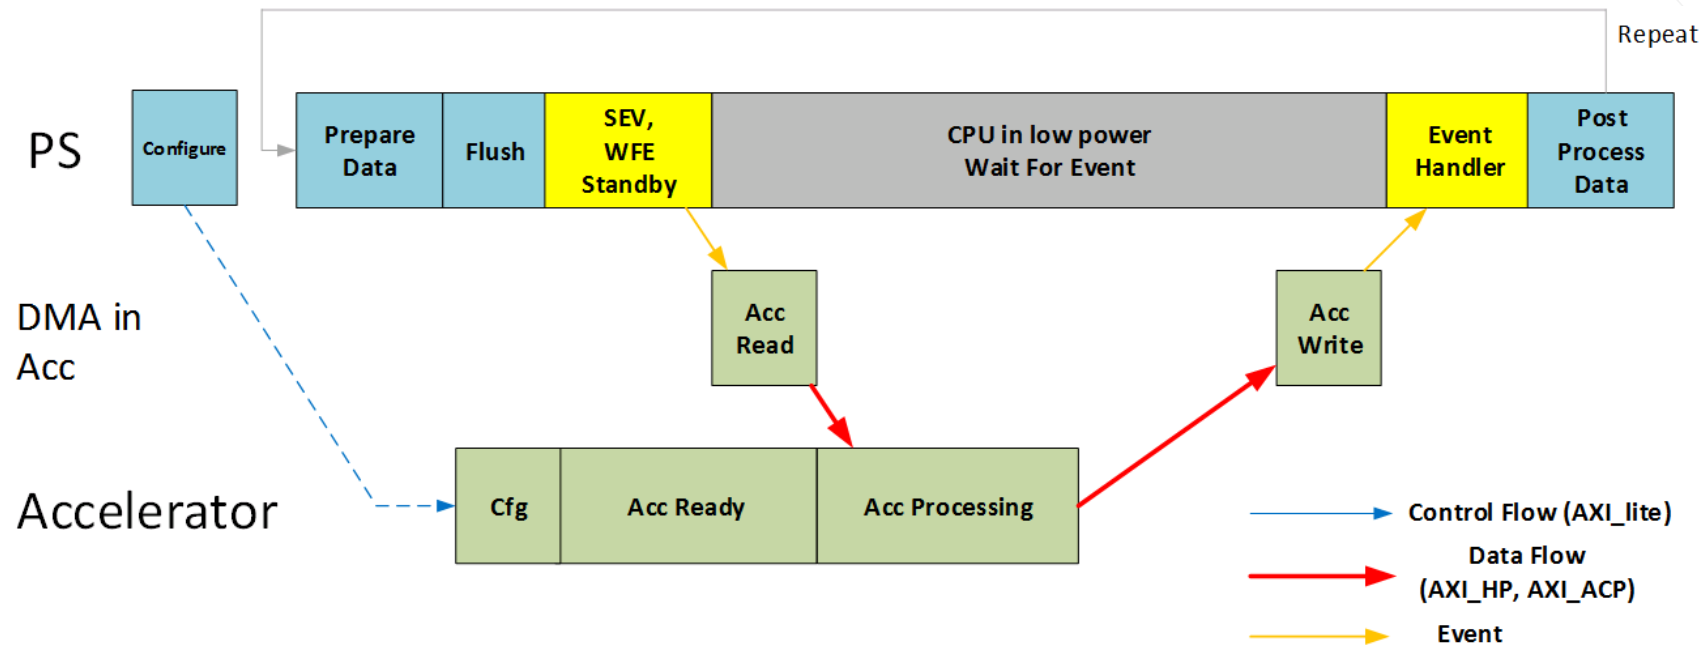
# Accelerator I/O Connected to PS: Non-Blocking




Flushing of Cache is required when the DMA can only access the DDR and not the Cache directly



# Accelerator I/O Connected to PS: Blocking for Power



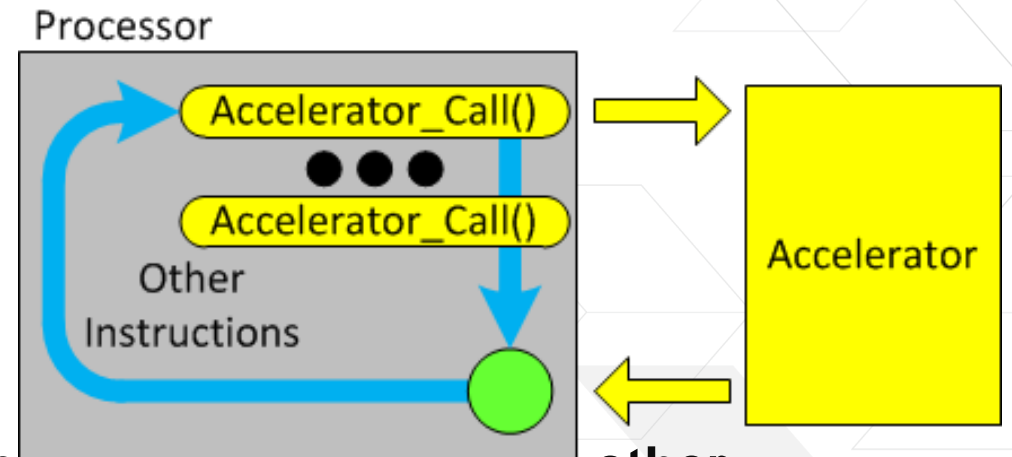
Flushing of Cache is required when the DMA can only access the DDR and not the Cache directly



# Synchronization and Pipelining Features

## > Consider code that makes several consecutive calls to the same accelerator function

- >> If each call is done asynchronously then other code can be run while the '*task*' of calling an accelerator is run
- >> If the *task* can be broken into independent *stages* there is opportunity for parallelism
- >> Capitalize on parallelism by restructuring code to create a *task pipeline*




## > The following *stages* of an accelerator task can be independent of each other

- >> Sending input data from shared memory to local accelerator buffers
- >> Processing data
- >> Returning results from local accelerator buffers to shared memory




# Pipelining Multiple Calls to the Accelerator

## > Sequential execution of matrix multiply calls (blocking)



## > Asynchronous pipelined calls *require* use of multi-buffers



**~7.5X speedup over software version**

# Synchronization of Software/Hardware Communication

- > **Synchronization is important so that the hardware and software communicate effectively**
  - >> Processor has to know when the accelerator is done
- > **Software determines if a hardware accelerator has completed processing data using one of two methods**
  - >> Interrupt driven (default) or polling
  - >> The desired method can be selected using an SDSoC tool-specific compiler flag
    - `--poll-mode <0/1>` where 1 = poll mode enabled
  - >> Contrary to common belief, polling can often yield better results for Linux-based systems
    - Due to overhead associated with handling interrupts in Linux

# Summary



# Summary

- > **C-Callable library consists of IP blocks written in VHDL or Verilog which can be called from user applications developed in SDSoC by statically linking it**
- > **C-Callable library function consists of two components**
  - >> Header file
  - >> Static library
- > **The SDSoC development environment offers a mechanism for generating multiple instances of an accelerator**
- > **#pragma SDS async(id) before the function call directs the SDSoC tool to return control after setting up all data transfers**
- > **#pragma SDS wait(id) must be used at suitable synchronization points**

# Summary

- > **The SDSoC development environment determines the accelerator-to-accelerator connections by analyzing the code that calls accelerators**
- > **Asynchronous/non-blocking accelerator calls allow accelerator tasks to be pipelined**
- > **A non-blocking accelerator call is just one of several mechanisms to control synchronization and/or pipelining in the SDSoC development environment**

**Adaptable.**  
**Intelligent.**

